

# CSC4005 High Performance Computing (2018/19)

## Assignment 2 - Parallel Searching using MPI

Contact: Dr Blesson Varghese  
*b.varghese@qub.ac.uk*

**Deadline: Wednesday, November 14, 23:59**

This assignment builds on the previous assignment, please re-read carefully the specification given in Assignment 1. The primary aim is to design, implement and experiment with MPI programs to perform the same search process. You will be provided with a single text file, `text.txt`, and eight different pattern files, `pattern1.txt`, `pattern2.txt`, ... , `pattern8.txt`. These can be found in the folder `inputs`.

### Test0

1. Modify the program `searching_sequential.c` so that it searches for the 8 patterns, sequentially, in a single batch run. You should use, `$time ./searching_sequential` to determine the total elapsed time which will be used as a reference point for comparison with your MPI implementations.

### Test1

1. Re-engineer `searching_sequential.c` into a coarse-grain embarrassingly parallel MPI program, `searching_MPI_0.c`. Here each MPI process will perform a sequential search for a single pattern. For example, with 2 MPI processes  $P_0$  will search for patterns 1, 3, 5 and 7 and  $P_1$  will search for patterns 2, 4, 6 and 8. With 4 MPI processes  $P_0$  will search for patterns 1 and 5,  $P_1$  will search for patterns 2 and 6,  $P_2$  will search for patterns 3 and 7 and  $P_3$  will search for patterns 4 and 8. With 8 MPI processes  $P_i$  will search for pattern  $i + 1$  (since you are only allocated 4 MPI slots on the Kelvin cluster, you will need to use `--oversubscribe`).
2. Using 4 cores in a batch job on the Kelvin cluster, run `searching_MPI_0.c` with 2, 4, and 8 processes. Use time to determine the elapsed time.
  - a. Draw a graph of the parallel speedup (PS) vs the number of processes. PS is the elapsed time taken by the sequential program divided by the elapsed time taken using P processes.
  - b. Draw a graph of the parallel efficiency (PE) vs the number of processes. PE is simply PS/P.

### Test2

Consider the results of Test0 and Test1.

1. How can you improve the PS of `searching_MPI_0.c` **without** modifying the program. Note that this improvement will be specific to this input data.
2. Verify your hypothesis by running `searching_MPI_0.c` with 2, 4 and 8 processes.
  - a. Draw a graph of the parallel speedup (PS) vs the number of processes.
  - b. Draw a graph of the parallel efficiency (PE) vs the number of processes.

### Test3

1. Design a fine-grain MPI program `searching_MPI_1.c`. This program should search for the 8 patterns in sequence but each search should be performed using  $n$  MPI processes, where  $n = 2, 4$  or  $8$ . You must adopt a master-slave model. First,  $P_0$ , the master process, should read the text file, `text.txt`, and send the appropriate part to the slave processes. For each of the 8 patterns in sequence the master:
  - a. should send the pattern to each slave;
  - b. all processes, including the master, should search their allocated portion of the text;
  - c. the master should collect the search results from the slave processes using `MPI_Reduce` and report the overall search result.
2. Using 4 cores in a batch job on the Kelvin Cluster, run `searching_MPI_1.c` with 2, 4, and 8 processes. Use time to determine the elapsed time.
  - a. Draw a graph of the parallel speedup (PS) vs the number of processes.
  - b. Draw a graph of the parallel efficiency (PE) vs the number of processes.
3. Comment on the performance of `searching_MPI_1.c`. Discuss how the performance might be improved.

### Submission

You should submit your assignment to Queen's Online as a zip file named `<student-number>.zip`. This must contain:

- A **source** folder, which contains the source code of the sequential program and two MPI programs along with corresponding job scripts.
- An **output** folder, which contains the output of executing your three programs and/or screenshots of the output.
- A **report** folder, which contains a high quality report addressing all points raised above. The content of the report must be organised in relevant sections.

Note: The program names should be as specified. There will be penalties for not adhering to the constraints provided in the brief. Do not include the input folder containing the test cases with your submission.

**All the best!**